

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
data = pd.read_csv("AB_NYC_2019.csv")
```

```
data.head()
```

```
↳
```

latitude	longitude	room_type	price	minimum_nights	number_of_reviews	last_review	r
40.64749	-73.97237	Private room	149	1	9	2018-10-19	
40.75362	-73.98377	Entire home/apt	225	1	45	2019-05-21	
40.80902	-73.94190	Private room	150	3	0	NaN	
40.68514	-73.95976	Entire home/apt	89	1	270	2019-07-05	
40.79851	-73.94399	Entire home/apt	80	10	9	2018-11-19	

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 48895 entries, 0 to 48894
Data columns (total 16 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                     48895 non-null  int64
1   name                                  48879 non-null  object
2   host_id                               48895 non-null  int64
3   host_name                             48874 non-null  object
4   neighbourhood_group                   48895 non-null  object
5   neighbourhood                         48895 non-null  object
6   latitude                             48895 non-null  float64
7   longitude                             48895 non-null  float64
8   room_type                             48895 non-null  object
9   price                                 48895 non-null  int64
10  minimum_nights                        48895 non-null  int64
11  number_of_reviews                     48895 non-null  int64
12  last_review                           38843 non-null  object
```

```

13 reviews_per_month      38843 non-null float64
14 calculated_host_listings_count  48895 non-null int64
15 availability_365         48895 non-null int64
dtypes: float64(3), int64(7), object(6)
memory usage: 6.0+ MB

```

```
from sklearn.preprocessing import LabelEncoder, OneHotEncoder
```

```
data['room_type'].value_counts()
```

```

Entire home/apt    25409
Private room       22326
Shared room        1160
Name: room_type, dtype: int64

```

Label Encoder

```

le = LabelEncoder()
data['room_type'] = le.fit_transform(data['room_type'])

```

```
data['room_type'].value_counts()
```

```

0    25409
1    22326
2     1160
Name: room_type, dtype: int64

```

```
le.classes_
```

```
array(['Entire home/apt', 'Private room', 'Shared room'], dtype=object)
```

One Hot Encoder

```
data['neighbourhood_group'].value_counts()
```

```

Manhattan    21661
Brooklyn     20104
Queens        5666
Bronx         1091
Staten Island   373
Name: neighbourhood_group, dtype: int64

```

```

oh = OneHotEncoder()
transformed_data = oh.fit_transform(data['neighbourhood_group'].values.reshape(-1,1)).toarray

```

```
oh.categories_
```

```
[array(['Bronx', 'Brooklyn', 'Manhattan', 'Queens', 'Staten Island'],
      dtype=object)]
```

```
transformed_data = pd.DataFrame(transformed_data, columns=['Bronx', 'Brooklyn', 'Manhattan', 'Staten Island', 'Queens'])
```

```
transformed_data.head()
```

	Bronx	Brooklyn	Manhattan	Queens	Staten Island
0	0.0	1.0	0.0	0.0	0.0
1	0.0	0.0	1.0	0.0	0.0
2	0.0	0.0	1.0	0.0	0.0
3	0.0	1.0	0.0	0.0	0.0
4	0.0	0.0	1.0	0.0	0.0

```
transformed_data.iloc[100,]
```

```
Bronx      0.0
Brooklyn    0.0
Manhattan   1.0
Queens      0.0
Staten Island 0.0
Name: 100, dtype: float64
```

```
data['neighbourhood_group'][100]
```

```
'Manhattan'
```

Normalization and Standardization

```
numeric_col = [c for c in data.columns if data[c].dtype != np.dtype('O')]
```

```
len(numeric_col), len(data.columns)
```

```
(11, 16)
```

```
numeric_col
```

```
['id',
 'host_id',
 'latitude',
 'longitude',
 'room_type',
 'price',
```

```
'minimum_nights',
'number_of_reviews',
'reviews_per_month',
'calculated_host_listings_count',
'availability_365']
```

```
temp = data[numeric_col]
```

```
temp
```

	id	host_id	latitude	longitude	room_type	price	minimum_nights	number_of_reviews
0	2539	2787	40.64749	-73.97237	1	149	1	
1	2595	2845	40.75362	-73.98377	0	225	1	
2	3647	4632	40.80902	-73.94190	1	150	3	
3	3831	4869	40.68514	-73.95976	0	89	1	
4	5022	7192	40.79851	-73.94399	0	80	10	
...
48890	36484665	8232441	40.67853	-73.94995	1	70	2	
48891	36485057	6570630	40.70184	-73.93317	1	40	4	
48892	36485431	23492952	40.81475	-73.94867	0	115	10	
48893	36485609	30985759	40.75751	-73.99112	2	55	1	
48894	36487245	68119814	40.76404	-73.98933	1	90	7	

48895 rows × 11 columns

Normalization

```
from sklearn.preprocessing import StandardScaler, MinMaxScaler
```

```
import warnings
warnings.filterwarnings('ignore')
```

```
normalizer = MinMaxScaler()
```

```
temp.dropna(axis = 1, inplace=True)
```

```
normalized_data = normalizer.fit_transform(temp)
```

```
pd.DataFrame(normalized_data,columns=temp.columns)
```

	id	host_id	latitude	longitude	room_type	price	minimum_nights	numbe
0	0.000000	0.000001	0.357393	0.511921	0.5	0.0149	0.000000	
1	0.000002	0.000001	0.614199	0.490469	0.0	0.0225	0.000000	
2	0.000030	0.000008	0.748252	0.569257	0.5	0.0150	0.001601	
3	0.000035	0.000009	0.448496	0.535649	0.0	0.0089	0.000000	
4	0.000068	0.000017	0.722820	0.565324	0.0	0.0080	0.007206	
...
48890	0.999929	0.030002	0.432502	0.554109	0.5	0.0070	0.000801	
48891	0.999940	0.023944	0.488906	0.585684	0.5	0.0040	0.002402	
48892	0.999950	0.085632	0.762117	0.556517	0.0	0.0115	0.007206	
48893	0.999955	0.112946	0.623612	0.476639	1.0	0.0055	0.000000	
48894	1.000000	0.248315	0.639412	0.480007	0.5	0.0090	0.004804	

48895 rows × 10 columns

Standardization

```
standard_scaler = StandardScaler()
```

```
s_data = standard_scaler.fit_transform(temp)
```

```
pd.DataFrame(s_data,columns=temp.columns)
```

	id	host_id	latitude	longitude	room_type	price	minimum_nights	n
0	-1.731277	-0.860159	-1.493849	-0.437652	0.909359	-0.015493	-0.293996	
1	-1.731272	-0.860158	0.452436	-0.684639	-0.924247	0.300974	-0.293996	
2	-1.731176	-0.860135	1.468399	0.222497	0.909359	-0.011329	-0.196484	
3	-1.731159	-0.860132	-0.803398	-0.164450	-0.924247	-0.265335	-0.293996	

Handling with missing values

```
data.isnull().sum()
```

```
id          0
name        16
host_id     0
host_name   21
neighbourhood_group  0
neighbourhood  0
latitude    0
longitude    0
room_type   0
price       0
minimum_nights  0
number_of_reviews  0
last_review 10052
reviews_per_month 10052
calculated_host_listings_count  0
availability_365  0
dtype: int64
```

Simple Imputer

```
from sklearn.impute import SimpleImputer
```

```
imputer = SimpleImputer(missing_values=np.nan , strategy='mean')
```

```
reviews_per_month_col = imputer.fit_transform(data['reviews_per_month'].values.reshape(-1,1))
```

```
pd.DataFrame(reviews_per_month_col).isnull().sum()
```

```
0    0
dtype: int64
```

```
data['reviews_per_month'].isnull().sum()
```

```
10052
```

Discretization

```
from sklearn.preprocessing import KBinsDiscretizer
```

```
temp.head()
```

ce	minimum_nights	number_of_reviews	calculated_host_listings_count	availability_365
49	1	9	6	365
25	1	45	2	355
50	3	0	1	365
89	1	270	1	194
80	10	9	1	0

Quantile Discretization Transform

```
trans = KBinsDiscretizer(n_bins=10,encode='ordinal',strategy='quantile')  
new_data = trans.fit_transform(temp)
```

```
pd.DataFrame(new_data,columns=temp.columns)
```

id host id latitude longitude room tvbe price minimum nights number of r

Double-click (or enter) to edit

Uniform Discretization Transform

```
trans = KBinsDiscretizer(n_bins=10, encode='ordinal', strategy='uniform')
new_data= trans.fit_transform(temp)
pd.DataFrame(new_data, columns=temp.columns)
```

	id	host_id	latitude	longitude	room_type	price	minimum_nights	number_of_r
0	0.0	0.0	3.0	5.0	5.0	0.0	0.0	
1	0.0	0.0	6.0	4.0	0.0	0.0	0.0	
2	0.0	0.0	7.0	5.0	5.0	0.0	0.0	
3	0.0	0.0	4.0	5.0	0.0	0.0	0.0	
4	0.0	0.0	7.0	5.0	0.0	0.0	0.0	
...
48890	9.0	0.0	4.0	5.0	5.0	0.0	0.0	
48891	9.0	0.0	4.0	5.0	5.0	0.0	0.0	
48892	9.0	0.0	7.0	5.0	0.0	0.0	0.0	
48893	9.0	1.0	6.0	4.0	9.0	0.0	0.0	
48894	9.0	2.0	6.0	4.0	5.0	0.0	0.0	

48895 rows × 10 columns

Quantile Discretization Transform

```
trans = KBinsDiscretizer(n_bins=10, encode='ordinal', strategy='kmeans')
new_data=trans.fit_transform(temp)
pd.DataFrame(new_data, columns=temp.columns)
```


	id	host_id	latitude	longitude	room_type	price	minimum_nights	number_of_reviews
0	0.0	0.0	1.0	4.0	1.0	0.0	0.0	
1	0.0	0.0	6.0	4.0	0.0	1.0	0.0	
2	0.0	0.0	8.0	5.0	1.0	0.0	1.0	
3	0.0	0.0	3.0	5.0	0.0	0.0	0.0	
4	0.0	0.0	7.0	5.0	0.0	0.0	2.0	
...
48890	9.0	0.0	2.0	5.0	1.0	0.0	1.0	
48891	9.0	0.0	2.0	6.0	1.0	0.0	1.0	
...
48893	9.0	1.0	6.0	3.0	2.0	0.0	0.0	
48894	9.0	2.0	6.0	4.0	1.0	0.0	2.0	

48895 rows × 10 columns